1 BIC ATAC-seq Pipeline (v2.0.0)

1.1 Introduction

The results are delivered in one folder, `atacSeq`, with four parts. `macs` holds the peaks called for each sample. `atlas` holds the merged peak atlas of the project and the raw count matrix over it, which is the input to any normalization or differential analysis. `bigwig` holds a normalized signal track per sample for a genome browser. `metrics` holds the QC metrics and plots. Projects with replicates have a fifth part, `diff`, with the differential accessibility analysis. How each file was made is described in `methods.md`.

In the names below, `<sample>` is the sample ID of the BAM and `<no>` is the project number (for example `18143`). Coordinates are 0-based with an exclusive end (BED convention) unless a file says otherwise. Peaks, the atlas and the signal tracks cover the autosomes and the X and Y chromosomes only; mitochondrial, unplaced and other contigs are left out.

1.2 Directory Structure

```
atacSeq
├── atlas
│   ├── macsPeaksMerged.saf
│   ├── peaks_raw_fcCounts.txt
│   └── peaks_raw_fcCounts.txt.summary
├── bigwig
│   └── <sample>_postProcess.shifted.10mNorm.bw
├── macs
│   └── <sample>_postProcess.shifted
│       ├── <sample>_postProcess.shifted_peaks.narrowPeak
│       ├── <sample>_postProcess.shifted_summits.bed
│       ├── <sample>_postProcess.shifted_peaks.xls
│       ├── <sample>_postProcess.shifted_peaks.blacklistFiltered.narrowPeak
│       └── <sample>_postProcess.shifted_summits.blacklistFiltered.bed
├── metrics
│   ├── Proj_<no>__postInsDistribution.pdf
│   ├── Proj_<no>__ATACSeqQC.pdf
│   ├── <sample>_postProcess___INS.txt
│   ├── <sample>_postProcess___INS.pdf
│   ├── <sample>_postProcess___MITO.txt
│   ├── <sample>_postProcess.tss_enrich.csv
│   ├── <sample>_postProcess.tss_enrich.pdf
│   └── <sample>_postProcess.large_tss_enrich.png
└── diff                        (projects with replicates only)
    ├── Proj_<no>__DiffPeaksEdgeR_V4.xlsx
    └── Proj_<no>__DiffPeaks_V4.pdf
```

The filtered BAM files (`<sample>_postProcess.bam`) are not delivered. If you need them, please contact us.

1.3 Peaks (`macs`)

One folder per sample with the output of MACS2 and two files derived from it. The folder and file names carry the suffix `_postProcess.shifted` because the peaks were called on the filtered, Tn5-shifted reads.

MACS2 output. These three files are MACS2's own, delivered unchanged:

- `*_peaks.narrowPeak`: every peak MACS2 called at $p < 0.01$, in ENCODE narrowPeak format (BED6+4). The columns are chromosome, start, end, peak name, score, strand (always `.`), fold enrichment at the summit, -log10 p-value, -log10 q-value, and the position of the summit relative to the peak start. A peak with several summits appears once per summit, with a letter suffix on its name (`_peak_12a`, `_peak_12b`).
- `*_summits.bed`: one line per summit, the single base of highest signal in each peak, with the peak name and its -log10 p-value.
- `*_peaks.xls`: the same peaks as a tab-delimited text table (not an Excel file) with 1-based coordinates, peak length, summit position, pileup, -log10 p-value, fold enrichment and -log10 q-value. The header lines record the MACS2 command and settings.

Filtered peaks. The same peaks after the ENCODE post-processing steps (see `methods.md`): at most 300,000 peaks per sample, the most significant by p-value, clipped to the chromosome ends, with every peak that overlaps the ENCODE blacklist removed and the score capped at 1000. Peaks keep their MACS2 names, so a peak can be looked up in the three files above.

- `*_peaks.blacklistFiltered.narrowPeak`: the filtered peaks, same columns as `*_peaks.narrowPeak`, sorted by p-value.
- `*_summits.blacklistFiltered.bed`: the summits of the filtered peaks.

Use the `blacklistFiltered` files for any per-sample peak analysis. They are the peaks the atlas was built from. The per-sample peak counts in these files are the ones to read against ENCODE's data standards.

1.4 Peak Atlas and Counts (`atlas`)

1.4.1 Atlas (`macsPeaksMerged.saf`)

The peak atlas: the filtered peaks of all samples pooled and merged, with peaks within 500 bp of each other joined into one region. The atlas is one set of regions shared by every sample, so that the samples can be compared region by region. It is in SAF format, as read by featureCounts, with no header line and five tab-separated columns:

Column	Contents
`GeneID`	region name, `Peak_1`, `Peak_2`, … in genome order
`Chr`	chromosome
`Start`	start, 1-based
`End`	end, inclusive
`Strand`	always `+`

Note that SAF coordinates are 1-based and inclusive, unlike the BED and narrowPeak files. To convert a region to BED, subtract 1 from `Start`.

Because the atlas pools and merges the peaks of all samples, its number of regions grows with the number of samples and is not a per-sample peak count.

1.4.2 Raw counts (`peaks_raw_fcCounts.txt`)

The number of fragments of each sample in each atlas region, from featureCounts. The first line is a comment with the featureCounts command; the second is the header. The columns are the five atlas columns (`Geneid`, `Chr`, `Start`, `End`, `Strand`), the region `Length`, and then one column per sample, named by the path of its filtered BAM (`out/<sample>/<sample>_postProcess.bam`).

A fragment is a read pair; it is counted once in every region it overlaps by at least one base, so a fragment that spans two adjacent regions is counted in both. The counts are raw, not normalized. They are the input to the DESeq2 scale factors, the QC plots and the differential analysis, and are what to use for any analysis that needs counts.

1.4.3 Count summary (`peaks_raw_fcCounts.txt.summary`)

For each sample, the number of fragments assigned to the atlas (`Assigned`) and the number not assigned, by reason (the `Unassigned_*` rows; `Unassigned_NoFeatures` is fragments that overlap no atlas region). The sum of the column is the number of fragments in the filtered BAM.

1.5 Signal Tracks (`bigwig`)

`<sample>_postProcess.shifted.10mNorm.bw` is the coverage of the filtered, Tn5-shifted reads of the sample, scaled to 10 million reads (`10mNorm`), in bigWig format. The tracks can be loaded directly into IGV or the UCSC Genome Browser. Because every sample is scaled to the same number of reads, the tracks can be compared by eye across samples. The value at a base is the number of shifted reads covering it, multiplied by 10 million and divided by the sample's total number of shifted reads.

1.6 Metrics (`metrics`)

1.6.1 QC report (`Proj_<no>__ATACSeqQC.pdf`)

Four pages built from the count matrix, with samples colored by group:

1. The percentage of each sample's fragments that fall in the atlas (`InPeaks`) against those outside it. This is a signal-to-noise measure; a sample with a much lower percentage than the others may have failed. It is computed against the pooled atlas, so it is higher than ENCODE's per-sample FRiP and should not be read against ENCODE's FRiP standard.
2. The same as counts, in millions of fragments.
3. A principal component analysis of the samples on the normalized log2 counts over the atlas, with sample labels. Replicates of a group should lie together; samples that separate from their group on PC1 or PC2 are worth a look.
4. The same PCA without labels.

1.6.2 Insert-size distributions (`Proj_<no>__postInsDistribution.pdf`)

The insert-size (fragment-length) distribution of each sample's filtered reads, from Picard CollectInsertSizeMetrics:

1. All samples on one plot, density against insert size (0 to 1000 bp), colored by group with one line type per replicate.
2. The same on a log scale, which shows the nucleosomal peaks.
3. One panel per sample, six per page.

A successful ATAC-seq library shows a large nucleosome-free peak below about 100 bp and smaller peaks at about 200 bp (one nucleosome) and 400 bp (two nucleosomes).

The histogram covers 0 to 1000 bp. Its smallest insert is 31 bp: pairs with an insert of 30 bp or less are removed by the read filter (see the methods document).

The per-sample files behind the plot are `<sample>_postProcess___INS.txt`, the Picard metrics file (summary statistics such as the median and mode insert size, followed by the histogram as `insert_size` and count columns), and `<sample>_postProcess___INS.pdf`, Picard's own histogram plot.

1.6.3 Mitochondrial fraction (`<sample>_postProcess___MITO.txt`)

One line per sample with the columns `SAMPLE`, `MITO_CONTIG`, `READS` (all reads of the filtered BAM), `MITO_READS` (those on the mitochondrial contig) and `FRAC_MITO`. Mitochondrial reads are left out of the peaks, the atlas and the signal tracks. The fraction is computed after the read filter, which removes duplicates, so it is lower than a fraction computed on the raw alignments (as ENCODE reports it) and should not be read against that standard.

1.6.4 TSS enrichment (`<sample>_postProcess.tss_enrich.csv`)

The ENCODE TSS enrichment score, the standard measure of signal-to-noise for ATAC-seq: how much higher the signal is at transcription start sites than in the flanking regions 2 kb away. Columns:

Column	Contents
`sample_id`	sample
`tss_enrich`	the TSS enrichment score
`peak_pos_bp`	position of the maximum of the profile, relative to the TSS
`avg_noise`	mean signal in the flanking regions, the denominator of the score
`n_tss_used`, `n_tss_total`, `pct_tss_used`	TSSs scored, TSSs in the reference set, and their ratio; `pct_tss_used` should be close to 1
`n_mapped_reads`	reads of the filtered BAM
`read_len`	read length used for the computation
`bam_file`	the filtered BAM

ENCODE's reading of the score depends on the genome and on the TSS annotation. The scores here use ENCODE's GENCODE TSS sets (v24 for hg38, v19 for b37, vM7 for mm10); read the cutoffs as a guide:

Genome	Concerning	Acceptable	Ideal
hg38	< 5	5 to 7	> 7

Genome	Concerning	Acceptable	Ideal
b37	< 6	6 to 10	> 10
mm10	< 10	10 to 15	> 15

Two plots accompany the score: `<sample>_postProcess.tss_enrich.pdf`, the mean signal profile across the 4 kb window centered on the TSS, and `<sample>_postProcess.large_tss_enrich.png`, a heatmap of the signal at every TSS, sorted by signal, with the mean profile above it.

1.7 Differential Accessibility

For projects with replicates, pairwise comparisons between groups are delivered in the `diff` folder as two files, named with the project number and, when one was used, a run tag:

- `Proj_<no>__DiffPeaksEdgeR_V4.xlsx`: one `Summary` sheet with the number of significant regions per comparison, then one sheet per comparison, named `<group2>_v_<group1>`, listing the regions significant at FDR < 0.05. Columns: `PeakNo` (the atlas region), `logFC` (log2 fold change, group 2 over group 1: positive means more accessible in group 2), `logCPM` (mean log2 counts per million), `PValue`, `FDR` (Benjamini-Hochberg adjusted p-value), the region's `Chr`, `Start`, `End` (1-based, from the atlas), `Strand` and `Length`, and the ChIPseeker annotation: `SYMBOL` and `GENENAME` of the nearest gene, `annotation` (promoter, exon, intron, distal intergenic, …) and `distanceToTSS`.
- `Proj_<no>__DiffPeaks_V4.pdf`: the PCA of the samples (with and without labels), then for each comparison an MA plot (log fold change against mean log CPM) and a volcano plot (log fold change against p-value), with significant regions in red.

Samples whose group starts with `EXC` were excluded from the differential analysis.